

# hw-05\_sanjaybhargavsiddi

December 11, 2023

## 1 HW-05

**Author: Sanjay Bhargav Siddi** This document shows a comprehensive regression analysis on a the Life Expectancy dataset from the 2023 #tidytuesday repository. The process includes data exploration, preprocessing, Ordinary Least Squares (OLS) regression with assumption checks, and the implementation and evaluation of alternative regression methods such as Random Forests and SVR.

### 1.1 1. Data selection and exploration

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split, KFold, cross_val_score, GridSearchCV
import statsmodels.api as sm
from sklearn.utils import resample
from sklearn.linear_model import Ridge, Lasso

warnings.filterwarnings("ignore")
```

Getting the dataset from #tidytuesday

```
[2]: life_expectancy = pd.read_csv('https://raw.githubusercontent.com/
    rfordata/tidytuesday/master/data/2023/2023-12-05/life_expectancy.csv')
life_expectancy_different_ages = pd.read_csv('https://raw.githubusercontent.com/
    rfordata/tidytuesday/master/data/2023/2023-12-05/
    life_expectancy_different_ages.csv')
life_expectancy_female_male = pd.read_csv('https://raw.githubusercontent.com/
    rfordata/tidytuesday/master/data/2023/2023-12-05/
    life_expectancy_female_male.csv')
```

```
[3]: merged_data = pd.merge(life_expectancy, life_expectancy_different_ages,
    ↪on=['Entity', 'Code', 'Year'])
data = pd.merge(merged_data, life_expectancy_female_male, on=['Entity', 'Code',
    ↪'Year'])
```

```
[4]: data
```

```
[4]:
```

	Entity	Code	Year	LifeExpectancy	LifeExpectancy0	\
0	Afghanistan	AFG	1950	27.7275	27.7275	
1	Afghanistan	AFG	1951	27.9634	27.9634	
2	Afghanistan	AFG	1952	28.4456	28.4456	
3	Afghanistan	AFG	1953	28.9304	28.9304	
4	Afghanistan	AFG	1954	29.2258	29.2258	
...	...	...	...	...	...	
19485	Zimbabwe	ZWE	2017	60.7095	60.7095	
19486	Zimbabwe	ZWE	2018	61.4141	61.4141	
19487	Zimbabwe	ZWE	2019	61.2925	61.2925	
19488	Zimbabwe	ZWE	2020	61.1242	61.1242	
19489	Zimbabwe	ZWE	2021	59.2531	59.2531	

	LifeExpectancy10	LifeExpectancy25	LifeExpectancy45	LifeExpectancy65	\
0	49.1459	54.442200	63.422500	73.4901	
1	49.2941	54.564400	63.500603	73.5289	
2	49.5822	54.799800	63.647600	73.6018	
3	49.8634	55.028603	63.788902	73.6706	
4	49.9306	55.116500	63.848100	73.7041	
...	...	...	...	...	
19485	64.6277	66.110596	71.014100	78.5895	
19486	65.1821	66.604500	71.267200	78.6681	
19487	65.0582	66.491600	71.203400	78.6739	
19488	64.8006	66.086900	70.519104	78.0986	
19489	62.8058	64.169700	68.801200	76.8507	

	LifeExpectancy80	LifeExpectancyDiffFM
0	83.7259	1.261900
1	83.7448	1.270601
2	83.7796	1.288300
3	83.8118	1.306601
4	83.8334	1.276501
...	...	...
19485	86.8135	4.748299
19486	86.8399	4.625503
19487	86.8614	5.017799
19488	86.5717	5.732201
19489	85.7716	5.812798

```
[19490 rows x 11 columns]
```

## Dataset Description:

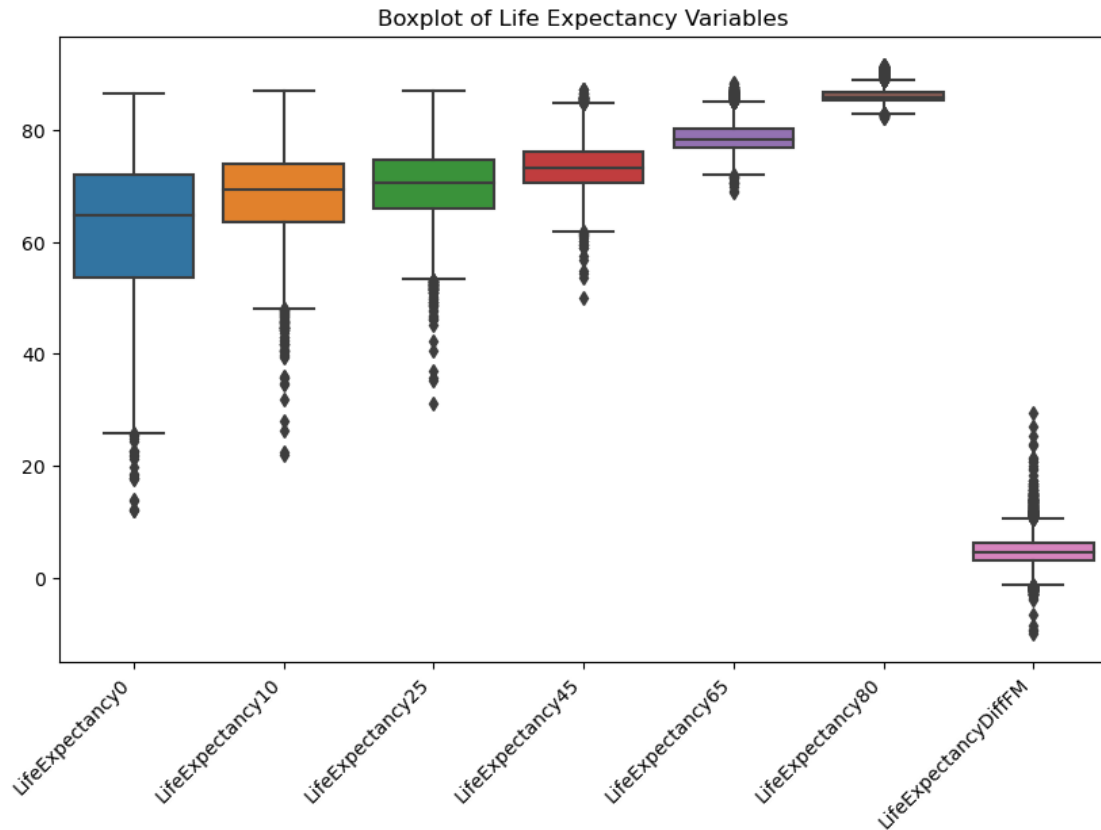
Variable	Class	Description
Entity	character	Country or region entity
Code	character	Entity code
Year	double	Year
LifeExpectancy0	double	Period life expectancy at birth - Sex: all - Age: 0
LifeExpectancy10	double	Period life expectancy - Sex: all - Age: 10
LifeExpectancy25	double	Period life expectancy - Sex: all - Age: 25
LifeExpectancy45	double	Period life expectancy - Sex: all - Age: 45
LifeExpectancy65	double	Period life expectancy - Sex: all - Age: 65
LifeExpectancy80	double	Period life expectancy - Sex: all - Age: 80
LifeExpectancyDiffFM	double	Life expectancy difference (f-m) - Type: period - Sex: both - Age: 0

```
[5]: # Describe missing values
missing_values = data.isnull().sum()
missing_values
```

```
[5]: Entity          0
Code          1168
Year          0
LifeExpectancy 0
LifeExpectancy0 0
LifeExpectancy10 0
LifeExpectancy25 0
LifeExpectancy45 0
LifeExpectancy65 0
LifeExpectancy80 0
LifeExpectancyDiffFM 0
dtype: int64
```

There are 1168 rows which have null values in the Code column.

```
[6]: # Plot outliers
plt.figure(figsize=(10, 6))
ax = sns.boxplot(data=data[['LifeExpectancy0', 'LifeExpectancy10',
↪ 'LifeExpectancy25', 'LifeExpectancy45', 'LifeExpectancy65',
↪ 'LifeExpectancy80', 'LifeExpectancyDiffFM']])
plt.title('Boxplot of Life Expectancy Variables')
ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha='right')
plt.show()
```



From this box plot we can easily identify the outliers. We can observe the presence of extreme values in the dataset by looking at the data points below the whiskers.

Example 1: In LifeExpectancy0 and LifeExpectancy1, the median line is not in the center suggesting that the data is not symmetrically distributed. The tall box size of LifeExpectancy0 indicates the narrow spread of data in it.

Example 2: In LifeExpectancy45, the outliers are present at either ends of the whiskers. And the wide nature of the box suggests a wider spread.

```
[7]: data_shape = data.shape
     data_shape
```

```
[7]: (19490, 11)
```

```
[8]: correlation_matrix = merged_data.corr()
     correlation_matrix
```

```
[8]:
```

	Year	LifeExpectancy	LifeExpectancy0	LifeExpectancy10	\
Year	1.000000	0.597792	0.597792	0.537543	
LifeExpectancy	0.597792	1.000000	1.000000	0.968572	
LifeExpectancy0	0.597792	1.000000	1.000000	0.968572	

LifeExpectancy10	0.537543	0.968572	0.968572	1.000000
LifeExpectancy25	0.523513	0.953081	0.953081	0.996099
LifeExpectancy45	0.521685	0.917104	0.917104	0.968553
LifeExpectancy65	0.528763	0.876127	0.876127	0.919194
LifeExpectancy80	0.495097	0.796174	0.796174	0.828120

	LifeExpectancy25	LifeExpectancy45	LifeExpectancy65	\
Year	0.523513	0.521685	0.528763	
LifeExpectancy	0.953081	0.917104	0.876127	
LifeExpectancy0	0.953081	0.917104	0.876127	
LifeExpectancy10	0.996099	0.968553	0.919194	
LifeExpectancy25	1.000000	0.983986	0.937024	
LifeExpectancy45	0.983986	1.000000	0.977651	
LifeExpectancy65	0.937024	0.977651	1.000000	
LifeExpectancy80	0.844521	0.898823	0.961757	

	LifeExpectancy80
Year	0.495097
LifeExpectancy	0.796174
LifeExpectancy0	0.796174
LifeExpectancy10	0.828120
LifeExpectancy25	0.844521
LifeExpectancy45	0.898823
LifeExpectancy65	0.961757
LifeExpectancy80	1.000000

The correlation between the “Year” variable and “Life Expectancy” is 0.60, indicating a moderate positive correlation.

All variables related to life expectancy show strong positive correlations with each other, ranging from 0.80 to 1.00

There is a decreasing trend in the correlation as the age groups get farther apart. For example, “LifeExpectancy” and “LifeExpectancy80” have a lower correlation (0.50), indicating a weaker positive relationship compared to the correlations within closer age groups

---

## 1.2 2. Data Preprocessing

```
[9]: Q1 = data.quantile(0.25)
      Q3 = data.quantile(0.75)
      IQR = Q3 - Q1
      outliers = ((data < (Q1 - 1.5 * IQR)) | (data > (Q3 + 1.5 * IQR))).any(axis=1)

      cleaned_data = data[~outliers]
```

```
[10]: cleaned_data
```

```
[10]:
```

	Entity	Code	Year	LifeExpectancy	LifeExpectancy0	\
0	Afghanistan	AFG	1950	27.7275	27.7275	
1	Afghanistan	AFG	1951	27.9634	27.9634	
2	Afghanistan	AFG	1952	28.4456	28.4456	
3	Afghanistan	AFG	1953	28.9304	28.9304	
4	Afghanistan	AFG	1954	29.2258	29.2258	
...	...	...	...	...	...	
19485	Zimbabwe	ZWE	2017	60.7095	60.7095	
19486	Zimbabwe	ZWE	2018	61.4141	61.4141	
19487	Zimbabwe	ZWE	2019	61.2925	61.2925	
19488	Zimbabwe	ZWE	2020	61.1242	61.1242	
19489	Zimbabwe	ZWE	2021	59.2531	59.2531	

	LifeExpectancy10	LifeExpectancy25	LifeExpectancy45	LifeExpectancy65	\
0	49.1459	54.442200	63.422500	73.4901	
1	49.2941	54.564400	63.500603	73.5289	
2	49.5822	54.799800	63.647600	73.6018	
3	49.8634	55.028603	63.788902	73.6706	
4	49.9306	55.116500	63.848100	73.7041	
...	...	...	...	...	
19485	64.6277	66.110596	71.014100	78.5895	
19486	65.1821	66.604500	71.267200	78.6681	
19487	65.0582	66.491600	71.203400	78.6739	
19488	64.8006	66.086900	70.519104	78.0986	
19489	62.8058	64.169700	68.801200	76.8507	


  

	LifeExpectancy80	LifeExpectancyDiffFM
0	83.7259	1.261900
1	83.7448	1.270601
2	83.7796	1.288300
3	83.8118	1.306601
4	83.8334	1.276501
...	...	...
19485	86.8135	4.748299
19486	86.8399	4.625503
19487	86.8614	5.017799
19488	86.5717	5.732201
19489	85.7716	5.812798

[17437 rows x 11 columns]

```
[11]: # Plot outliers
plt.figure(figsize=(10, 6))
ax = sns.boxplot(data=cleaned_data[['LifeExpectancy0', 'LifeExpectancy10',
↳ 'LifeExpectancy25', 'LifeExpectancy45', 'LifeExpectancy65',
↳ 'LifeExpectancy80', 'LifeExpectancyDiffFM']])
plt.title('Boxplot of the cleaned data')
```

```
ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha='right')
plt.show()
```



```
[12]: df = cleaned_data
      #f = df[['Year', 'Entity', 'LifeExpectancy']]
```

### 1.2.1 Label Encoding: Transform categorical values into numerical labels

```
[13]: label_encoder = LabelEncoder()
      # Applying label encoding to "Entity" column
      df['Entity'] = label_encoder.fit_transform(df['Entity'])
```

## 1.3 3. Ordinary Least Squares (OLS) Regression

### 1.3.1 Can we predict life expectancy at age 45 based on the selected features?

```
[14]: X = df.drop(['LifeExpectancy', 'LifeExpectancy0', 'Code', 'LifeExpectancyDiffFM'],
      ↪axis=1) # Features
      y = df['LifeExpectancy0'] # Target variable
```

```
[15]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
# Resampling
X_train_resampled, y_train_resampled = resample(X_train, y_train, replace=True,
↳ random_state=42)
```

```
[16]: # Model Building with statsmodels
X_train_sm = sm.add_constant(X_train_resampled)

# Adding a constant for the intercept
ols_model = sm.OLS(y_train_resampled, X_train_sm).fit()

# Evaluating the Model Performance
# Adding a constant for the intercept in test data (statsmodels)
X_test_sm = sm.add_constant(X_test)

# Predicting using statsmodels model
y_pred_sm = ols_model.predict(X_test_sm)

# Model Diagnostics
ols_model.summary()
```

[16]:

<b>Dep. Variable:</b>	LifeExpectancy0	<b>R-squared:</b>	0.966
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.966
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	5.673e+04
<b>Date:</b>	Mon, 11 Dec 2023	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	19:20:00	<b>Log-Likelihood:</b>	-30292.
<b>No. Observations:</b>	13949	<b>AIC:</b>	6.060e+04
<b>Df Residuals:</b>	13941	<b>BIC:</b>	6.066e+04
<b>Df Model:</b>	7		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P>  t	[0.025	0.975]
const	-63.6313	3.410	-18.661	0.000	-70.315	-56.948
Entity	0.0001	0.000	0.508	0.611	-0.000	0.001
Year	0.0166	0.001	17.568	0.000	0.015	0.018
LifeExpectancy10	5.6042	0.066	84.448	0.000	5.474	5.734
LifeExpectancy25	-5.4611	0.126	-43.312	0.000	-5.708	-5.214
LifeExpectancy45	1.1135	0.120	9.315	0.000	0.879	1.348
LifeExpectancy65	-0.0165	0.116	-0.143	0.886	-0.243	0.210
LifeExpectancy80	0.1356	0.081	1.667	0.095	-0.024	0.295

<b>Omnibus:</b>	2140.528	<b>Durbin-Watson:</b>	1.985
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	20844.814
<b>Skew:</b>	-0.432	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	8.926	<b>Cond. No.</b>	3.78e+05

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



[2] The condition number is large,  $3.78e+05$ . This might indicate that there are strong multicollinearity or other numerical problems.

R-squared value is 96.6%

This suggests a good fit and indicates that approximately 96.6% of the variance in the dependent variable (LifeExpectancy0) is explained by the independent variables in the model

The p-values indicate the statistical significance of each predictor. A small p-value suggests that the variable is likely a meaningful addition to the model. The significant coefficients with low p-values suggest that 'Year,' 'LifeExpectancy10,' 'LifeExpectancy25,' 'LifeExpectancy45' are likely important predictors of 'LifeExpectancy0'.

The high condition number suggests potential multicollinearity, indicating that some predictors might be highly correlated

```
[17]: residuals = y_test - y_pred_sm
dw_test = sm.stats.stattools.durbin_watson(residuals)
print(f"Durbin-Watson test statistic: {dw_test}")
```

Durbin-Watson test statistic: 1.97398253794617

There is little evidence of significant autocorrelation in the residuals, but since it is within the 1.5 to 2.5 range we can say there is no significant autocorrelation.

### 1.3.2 OLS Regression with CV

```
[18]: # 5 Folds
kf = KFold(n_splits=5, shuffle=True, random_state=42)

# Lists to store evaluation metrics across folds
r_squared_values = []
mae_values = []

for train_index, val_index in kf.split(X_train_resampled):
    # Splitting data into train and validation sets for this fold
    X_train_fold, X_val_fold = X_train_resampled.iloc[train_index],
    ↪X_train_resampled.iloc[val_index]
    y_train_fold, y_val_fold = y_train_resampled.iloc[train_index],
    ↪y_train_resampled.iloc[val_index]

    # Fitting the OLS model
    ols_model_fold = sm.OLS(y_train_fold, sm.add_constant(X_train_fold)).fit()

    # Making predictions on the validation set
    y_pred_val = ols_model_fold.predict(sm.add_constant(X_val_fold))

    # Calculating R-squared for this fold
    r_squared_fold = 1 - (np.sum((y_val_fold - y_pred_val) ** 2) / np.
    ↪sum((y_val_fold - np.mean(y_val_fold)) ** 2))
```

```

r_squared_values.append(r_squared_fold)

# Calculating MAE for this fold
mae_fold = np.mean(np.abs(y_val_fold - y_pred_val))
mae_values.append(mae_fold)

# Calculating average metrics across all folds
avg_r_squared = np.mean(r_squared_values)
avg_mae = np.mean(mae_values)

# Printing average metrics
print("Average R-squared across folds:", avg_r_squared)
print("Average MAE across folds:", avg_mae)

```

Average R-squared across folds: 0.9659915752082947

Average MAE across folds: 1.4288174488293415

```
[19]: ols_model_fold.summary()
```

[19]:

<b>Dep. Variable:</b>	LifeExpectancy0	<b>R-squared:</b>	0.966
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.966
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	4.472e+04
<b>Date:</b>	Mon, 11 Dec 2023	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	19:20:00	<b>Log-Likelihood:</b>	-24225.
<b>No. Observations:</b>	11160	<b>AIC:</b>	4.847e+04
<b>Df Residuals:</b>	11152	<b>BIC:</b>	4.852e+04
<b>Df Model:</b>	7		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P>  t	[0.025	0.975]
<b>const</b>	-66.6259	3.795	-17.558	0.000	-74.064	-59.188
<b>Entity</b>	4.551e-05	0.000	0.165	0.869	-0.000	0.001
<b>Year</b>	0.0166	0.001	15.784	0.000	0.015	0.019
<b>LifeExpectancy10</b>	5.5720	0.073	76.052	0.000	5.428	5.716
<b>LifeExpectancy25</b>	-5.4118	0.139	-39.028	0.000	-5.684	-5.140
<b>LifeExpectancy45</b>	1.1146	0.132	8.472	0.000	0.857	1.373
<b>LifeExpectancy65</b>	-0.0704	0.128	-0.548	0.584	-0.322	0.181
<b>LifeExpectancy80</b>	0.2046	0.091	2.260	0.024	0.027	0.382

<b>Omnibus:</b>	1711.010	<b>Durbin-Watson:</b>	1.989
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	17347.023
<b>Skew:</b>	-0.414	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	9.051	<b>Cond. No.</b>	3.77e+05

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.77e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Comparing:

Both models have high R-squared values, suggesting a good fit to the data. Model 1 has more observations and a higher F-statistic, but Model 2 has a lower AIC, indicating a more economical model. The differences in coefficients are expected due to the different subsets of data used in training and testing.

## 1.4 4. Other Regression Techniques

Scaling the dataset

```
[20]: scaler = StandardScaler()
      X_train_scaled = scaler.fit_transform(X_train)
      X_test_scaled = scaler.transform(X_test)
```

### 1.4.1 Ridge Regression

Ridge Regression is a linear regression technique that addresses the issue of multicollinearity in multiple linear regression. Multicollinearity occurs when independent variables in a regression model are highly correlated, leading to instability and inflated standard errors in parameter estimates.

Ridge Regression is widely used in machine learning and statistics to improve the stability and reliability of linear regression models, particularly when dealing with high-dimensional datasets or correlated predictors

```
[21]: ridge_params = {'alpha': [0.1, 1.0, 5.0, 10.0]}
      ridge_model = Ridge()
      ridge_grid = GridSearchCV(ridge_model, ridge_params, cv=5,
      ↪scoring='neg_mean_squared_error')
      ridge_grid.fit(X_train_scaled, y_train)
```

```
[21]: GridSearchCV(cv=5, estimator=Ridge(),
                  param_grid={'alpha': [0.1, 1.0, 5.0, 10.0]},
                  scoring='neg_mean_squared_error')
```

```
[22]: # Evaluating Ridge Regression
      ridge_predictions = ridge_grid.predict(X_test_scaled)
      ridge_mse = mean_squared_error(y_test, ridge_predictions)
      print(f'Ridge Regression Mean Squared Error: {ridge_mse}')
      print("Best Ridge Hyperparameters:", ridge_grid.best_params_)
```

Ridge Regression Mean Squared Error: 4.716492458393628

Best Ridge Hyperparameters: {'alpha': 0.1}

### 1.4.2 Lasso Regression

Lasso Regression is a linear regression technique used for variable selection and regularization. Similar to Ridge Regression, Lasso Regression aims to prevent overfitting and improve model interpretability by adding a penalty term to the ordinary least squares objective function.

```
[23]: lasso_params = {'alpha': [0.1, 1.0, 5.0, 10.0]}
      lasso_model = Lasso()
```

```
lasso_grid = GridSearchCV(lasso_model, lasso_params, cv=5,
    ↪scoring='neg_mean_squared_error')
lasso_grid.fit(X_train_scaled, y_train)
```

```
[23]: GridSearchCV(cv=5, estimator=Lasso(),
    param_grid={'alpha': [0.1, 1.0, 5.0, 10.0]},
    scoring='neg_mean_squared_error')
```

```
[24]: # Evaluating Lasso Regression
lasso_predictions = lasso_grid.predict(X_test_scaled)
lasso_mse = mean_squared_error(y_test, lasso_predictions)
print(f'Lasso Regression Mean Squared Error: {lasso_mse}')
print("Best Lasso Hyperparameters:", lasso_grid.best_params_)
```

Lasso Regression Mean Squared Error: 6.8459625667892885  
 Best Lasso Hyperparameters: {'alpha': 0.1}

```
[25]: # Interpretation of sparse coefficients in Lasso Regression
lasso_coefficients = lasso_grid.best_estimator_.coef_
print("\nLasso Regression Coefficients:")
for feature, coefficient in zip(X.columns, lasso_coefficients):
    print(f'{feature}: {coefficient}')
```

Lasso Regression Coefficients:  
 Entity: 0.04470436408778068  
 Year: 0.6219651203781972  
 LifeExpectancy10: 12.278419476747326  
 LifeExpectancy25: -0.0  
 LifeExpectancy45: -1.9369881134869404  
 LifeExpectancy65: -0.0  
 LifeExpectancy80: 0.2925450119403004

Entity, Year, LifeExpectancy10, and LifeExpectancy80 - The coefficients for these variables are positive, indicating that as the values increase, the predicted 'LifeExpectancy0' also increases.

LifeExpectancy25, LifeExpectancy65: 0.0 - The coefficients for 'LifeExpectancy25', and 'LifeExpectancy65' are zero, implying that these variables do not contribute to the prediction of 'LifeExpectancy0' in the model.

LifeExpectancy45: -1.9370 - The 'LifeExpectancy45' variable has a negative coefficient, suggesting a negative impact on the predicted 'LifeExpectancy0'.

### 1.4.3 Conclusion:

In terms of model performance, the OLS model has the lowest average MAE across folds, indicating better predictive accuracy compared to both Lasso and Ridge Regression. Ridge Regression has a lower Mean Squared Error (4.716) compared to Lasso Regression (6.846). This suggests that Ridge Regression is performing better in terms of overall prediction accuracy on the test data.

In terms of feature selection, lasso Regression tends to set some coefficients to exactly zero, leading to sparsity in the model. The zero coefficients imply that certain features are not contributing significantly to the predictions.

In summary, based on the provided results, OLS performs well in terms of prediction accuracy, while Ridge Regression outperforms Lasso Regression in terms of Mean Squared Error.